

Sequential Monte Carlo methods

Lecture 4 – The bootstrap particle filter

Fredrik Lindsten, Uppsala University
2017-08-24

Outline – Lecture 4

Aim: Derive our first sequential Monte Carlo method: the bootstrap particle filter.

Outline:

1. A (hopefully) intuitive preview
2. The bootstrap particle filter
3. Resampling
4. A toy example and a real world application


1/20

Particle filter preview

A (hopefully) intuitive preview (I/III)

Consider a toy 1D localization problem.

Data



Model

Dynamics:

$$X_{t+1} = X_t + u_t + V_t,$$

where X_t denotes position, u_t denotes velocity (known), $V_t \sim \mathcal{N}(0, 5)$ denotes an unknown disturbance.

Measurements:

$$Y_t = h(X_t) + E_t.$$

where $h(\cdot)$ denotes the world model (here the terrain height) and $E_t \sim \mathcal{N}(0, 1)$ denotes an unknown disturbance.

Task: Learn about the state X_t (position) based on the measurements $y_{1:t}$ by computing the filter density $p(x_t | y_{1:t})$.

2/20

Highlights two **key capabilities** of the PF:

1. Automatically handles an unknown and dynamically changing number of hypotheses (modes).
2. Works with nonlinear/non-Gaussian models.

The bootstrap particle filter

Nonlinear filtering problem

Recall that the nonlinear filtering problem amounts to computing the filter PDF $p(\mathbf{x}_t | y_{1:t})$ when the model is given by

$$\mathbf{X}_{t+1} | (\mathbf{X}_t = \mathbf{x}_t) \sim p(\mathbf{x}_{t+1} | \mathbf{x}_t),$$

$$Y_t | (\mathbf{X}_t = \mathbf{x}_t) \sim p(y_t | \mathbf{x}_t),$$

$$\mathbf{X}_0 \sim p(\mathbf{x}_0).$$

We have shown that the solution is

$$p(\mathbf{x}_t | y_{1:t}) = \frac{p(y_t | \mathbf{x}_t)p(\mathbf{x}_t | y_{1:t-1})}{p(y_t | y_{1:t-1})},$$

$$p(\mathbf{x}_t | y_{1:t-1}) = \int p(\mathbf{x}_t | \mathbf{x}_{t-1})p(\mathbf{x}_{t-1} | y_{1:t-1})d\mathbf{x}_{t-1}.$$

Basic idea: Try to approximate $p(\mathbf{x}_t | y_{1:t})$ sequentially in time $t = 0, 1, \dots$ using importance sampling!

Particle filter – representation

The **particle filter** approximates $p(\mathbf{x}_t | y_{1:t})$ by maintaining an **empirical distribution** made up of N samples (particles) $\{\mathbf{x}_t^i\}_{i=1}^N$ and corresponding importance weights $\{w_t^i\}_{i=1}^N$

$$\underbrace{\hat{p}^N(\mathbf{x}_t | y_{1:t})}_{\hat{\pi}^N(\mathbf{x}_t)} = \sum_{i=1}^N w_t^i \delta_{\mathbf{x}_t^i}(\mathbf{x}_t).$$

The particle filter provides a well-founded way of exploring the state space using random simulation.

6/20

Importance sampling reminder

Algorithm 1 Importance sampler

1. Sample $\mathbf{x}^i \sim q(\mathbf{x})$.
2. Compute the weights $\tilde{w}^i = \tilde{\pi}(\mathbf{x}^i)/q(\mathbf{x}^i)$.
3. Normalize the weights $w^i = \tilde{w}^i / \sum_{j=1}^N \tilde{w}^j$.

Each step is carried out for $i = 1, \dots, N$.

7/20

Sampling from the proposal

We sample from the proposal

$$q(\mathbf{x}_t | y_{1:t}) = \sum_{i=1}^N \nu_{t-1}^i q(\mathbf{x}_t | \mathbf{x}_{t-1}^i, y_t)$$

using a two step procedure:

1. Select one of the components

$$\mathbf{a}_t^i \sim \mathcal{C}(\{\nu_{t-1}^j\}_{j=1}^N) \quad (\text{categorical distribution})$$

2. Generate a sample from the selected component,

$$\mathbf{x}_t^i \sim q(\mathbf{x}_t | \mathbf{x}_{t-1}^{\mathbf{a}_t^i}, y_t)$$

Repeat this N times, for $i = 1, \dots, N$.

8/20

Selecting the mixture components – resampling

The particle $\bar{\mathbf{x}}_{t-1}^i = \mathbf{x}_{t-1}^{\mathbf{a}_t^i}$ is referred to as the **ancestor** of \mathbf{x}_t^i , since \mathbf{x}_t^i is generated conditionally on $\bar{\mathbf{x}}_{t-1}^i$.

The variable $\mathbf{a}_t^i \in \{1, \dots, N\}$ is referred to as the **ancestor index**, since it indexes the ancestor of particle \mathbf{x}_t^i at time $t - 1$.

Sampling the N ancestor indices

$$\mathbf{a}_t^i \sim \mathcal{C}(\{\nu_{t-1}^j\}_{j=1}^N), \quad i = 1, \dots, N$$

is referred to as **resampling**.

Resampling generates a new set of particles $\{\bar{\mathbf{x}}_{t-1}^i\}_{i=1}^N$ by **sampling with replacement** from among $\{\mathbf{x}_{t-1}^j\}_{j=1}^N$, according to some weights $\{\nu_{t-1}^j\}_{j=1}^N$.

9/20

Next step – computing the weights

Algorithm 2 Importance sampler

1. Sample $x^i \sim q(x)$.
2. **Compute the weights** $\tilde{w}^i = \tilde{\pi}(x^i)/q(x^i)$.
3. Normalize the weights $w^i = \tilde{w}^i / \sum_{j=1}^N \tilde{w}^j$.

Each step is carried out for $i = 1, \dots, N$.

10/20

Result – A first particle filter

Algorithm 3 Bootstrap particle filter (for $i = 1, \dots, N$)

1. **Initialization** ($t = 0$):
 - (a) Sample $x_0^i \sim p(x_0)$.
 - (b) Set initial weights: $w_0^i = 1/N$.
2. **for** $t = 1$ **to** T **do**
 - (a) **Resample**: sample ancestor indices $a_t^i \sim \mathcal{C}(\{w_{t-1}^j\}_{j=1}^N)$.
 - (b) **Propagate**: sample $x_t^i \sim p(x_t | x_{t-1}^{a_t^i})$.
 - (c) **Weight**: compute $\tilde{w}_t^i = p(y_t | x_t^i)$ and normalize $w_t^i = \tilde{w}_t^i / \sum_{j=1}^N \tilde{w}_t^j$.

11/20

SMC structure



Same structure for all SMC algorithms.

For the bootstrap PF, given $\{x_{t-1}^i, w_{t-1}^i\}_{i=1}^N$:

Resampling: $a_t^i \sim \mathcal{C}(\{w_{t-1}^j\}_{j=1}^N)$.

Propagation: $x_t^i \sim p(x_t | x_{t-1}^{a_t^i})$.

Weighting: $\tilde{w}_t^i = p(y_t | x_t^i)$ and normalize.

The result is a new weighted set of particles $\{x_t^i, w_t^i\}_{i=1}^N$.

12/20

Intermediate approximations

Approximation of filtering distribution at time $t - 1$:

$$\sum_{i=1}^N w_{t-1}^i \delta_{x_{t-1}^i}(x_{t-1}) \approx p(x_{t-1} | y_{1:t-1}).$$

For the **bootstrap particle filter**:

- After resampling: $\frac{1}{N} \sum_{i=1}^N \delta_{\bar{x}_{t-1}^i}(x_{t-1}) \approx p(x_{t-1} | y_{1:t-1})$.
- After propagation: $\frac{1}{N} \sum_{i=1}^N \delta_{x_t^i}(x_t) \approx p(x_t | y_{1:t-1})$.
- After weighting: $\sum_{i=1}^N w_t^i \delta_{x_t^i}(x_t) \approx p(x_t | y_{1:t})$.

13/20

Examples

An LG-SSM example (I/II)

Whenever you are working on a nonlinear inference method, always make sure that it solves the linear special case first!

Consider the following LG-SSM (simple 1D positioning example)



$$\begin{pmatrix} X_t^{\text{pos}} \\ X_t^{\text{vel}} \\ X_t^{\text{acc}} \end{pmatrix} = \begin{pmatrix} 1 & T_s & T_s^2/2 \\ 0 & 1 & T_s \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} X_{t-1}^{\text{pos}} \\ X_{t-1}^{\text{vel}} \\ X_{t-1}^{\text{acc}} \end{pmatrix} + \begin{pmatrix} T_s^3/6 \\ T_s^2/2 \\ T_s \end{pmatrix} V_t, \quad V_t \sim \mathcal{N}(0, Q),$$

$$Y_t = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} X_t^{\text{pos}} \\ X_t^{\text{vel}} \\ X_t^{\text{acc}} \end{pmatrix} + E_t, \quad E_t \sim \mathcal{N}(0, R).$$



The Kalman filter provides the true filtering density, which implies that we can compare the PF to the truth in this case.

14/20

An LG-SSM example (II/II)



Using 200 particles.



Using 20 000 particles.

The particle filter estimates converge as the number of particles tends to infinity (Lecture 5).


15/20

Nonlinear real-world application example

Aim: Compute the **position** of a person moving around indoors using sensors (inertial, magnetometer and radio) located in an ID badge, and a map.




The sensors (IMU and radio) and the DSP are mounted inside an ID badge.




The inside of the ID badge.

16/20

Application – indoor localization (II/III)



“Likelihood model” for an office environment, the bright areas are rooms and corridors (i.e., walkable space).



An estimated trajectory and the particle cloud visualized at a particular instance.

17/20

Application – indoor localization (III/III)



[Show movie](#)

 Johan Kihlberg, Simon Tegelid, Manon Kok and Thomas B. Schön. **Map aided indoor positioning using particle filters.** *Reglermöte (Swedish Control Conference)*, Linköping, Sweden, June 2014.

18/20

Use of random numbers in the particle filter

Random numbers are used to

1. **initialize**
2. **resample** and
3. **propagate**

the particles.

The weighting step does not require any new random numbers, it is just a function of already existing random numbers.

We can reason about and make use of the **joint probability distribution of these random variables**, from which the particle filter **generates one realization each time it is executed**.

19/20

A few concepts to summarize lecture 4

Bootstrap particle filter: A particle filter with a specific choice of proposals. Particles are simulated according to the dynamical model and weights are assigned according to the measurement likelihood.

Resampling: The procedure that generates a new set of particles $\{\bar{x}_{t-1}^i\}_{i=1}^N$ by sampling with replacement from among $\{x_{t-1}^j\}_{j=1}^N$, according to some weights $\{\nu_{t-1}^j\}_{j=1}^N$.

Ancestor indices: Random variable that are used to make the stochasticity of the resampling step explicit by keeping track of which particles that get resampled.

20/20