

# Sequential Monte Carlo methods

## Lecture 13 – Gibbs sampling

---

Fredrik Lindsten, Uppsala University  
2017-08-28

## Outline – Lecture 13

**Aim:** Show an alternative MCMC procedure (Gibbs sampling) and how it conceptually can be used for learning of dynamical systems

### Outline:

1. The Gibbs sampler
2. Gibbs sampling for linear dynamical systems
3. Gibbs sampler for nonlinear dynamical systems

1/19

## The Gibbs sampler

---

## Challenge with MCMC

Designing **efficient** Metropolis–Hastings kernels for **arbitrary and high-dimensional** target distributions can be very challenging.

Gibbs sampling turns the overall sampling problem into a **series of sub-problems**, each of which is hopefully easier to address.

2/19

## A first Gibbs sampler

Let  $\pi(x_1, x_2)$  be a target distribution over two (groups of) variables.

**Basic factorization:**  $\pi(x_1, x_2) = \pi(x_2 | x_1)\pi(x_1)$

**Thus:**

- If  $(X_1, X_2) \sim \pi(x_1, x_2)$ , then  $X_1$  is distributed according to  $\pi(x_1)$ .
- If  $X_2^* | (X_1 = x_1) \sim \pi(x_2 | x_1)$ , then  $(X_1, X_2^*)$  is distributed according to  $\pi(x_1, x_2)$ .

Starting with a sample from the joint distribution, we can replace any of the variables by a draw from it's full conditional and still have a sample from the joint distribution.

3/19

## ex) Gibbs sampler illustration

**ex)** Sample from,

$$\pi(x_1, x_2) = \mathcal{N} \left( \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \mid \begin{pmatrix} 10 \\ 10 \end{pmatrix}, \begin{pmatrix} 2 & 1 \\ 1 & 1 \end{pmatrix} \right).$$

**Gibbs sampler:**

**Initialize**  $x_1[1] = 0, x_2[1] = 0$

**for**  $m = 2, \dots, M$

    Draw  $x_1[m] \sim \pi(x_1 | x_2[m-1])$ ;

    Draw  $x_2[m] \sim \pi(x_2 | x_1[m])$ .

4/19

## ex) Gibbs sampler illustration

5/19

## MCMC kernels

An MCMC sampler generates the Markov chain  $\{x[m]\}_{m=1}^M$  by:

- **Initialize:** set  $x[1]$  arbitrarily.
- **For**  $m = 2$  **to**  $M$ : sample  $x[m] \sim \kappa(x[m-1], x^*)$ .

$\kappa(x, x^*)$  is a **Markov kernel** on  $\mathcal{X}$ , i.e. a conditional distribution for the next state  $x^*$  given the current state  $x$ .

**Basic requirement 1:** Stationarity of  $\pi(x)$ ,

$$\int \pi(x) \kappa(x, x^*) dx = \pi(x^*).$$

**Basic requirement 2:** Ergodicity —  $\kappa$  must allow the state to move in order to explore the state space.

6/19

## The Gibbs Markov kernel

**Target:**  $\pi(x) = \pi(x_1, \dots, x_d)$

**Input** a configuration  $x = (x_1, \dots, x_d)$

**for**  $j = 1, \dots, d$

    Sample  $x_j^* \sim \pi(x_j | x_1^*, \dots, x_{j-1}^*, x_{j+1}, \dots, x_d)$

**Output**  $x^* = (x_1^*, \dots, x_d^*)$ .

**Gibbs kernel:** This procedure defines a Markov kernel  $\kappa(x, x^*)$  with stationary distribution  $\pi(x)$ .

7/19

## Extensions

There are many possible extensions of the basic Gibbs procedure, which also result in valid MCMC kernels.

- **Random scan:** select components to sample randomly (with or without replacement)
- **Overlapping blocks:** the groups of variables need not be disjoint
- **Collapsing:** analytical marginalization of some of the variables (!)

8/19

## Extensions – Composition of MCMC methods

In many cases, exact sampling from some of the full conditional distributions is not possible.

Sufficient to sample from some Markov kernel which has the **full conditional distribution** as stationary distribution — we can make use of a combination of MCMC techniques.

If exact sampling from  $\pi(x_j | x_{-j})$  is not possible:

$$x_j^* \sim \kappa_j(x, x_j^*) \text{ where } \int \kappa_j(x, x_j^*) \pi(x_j | x_{-j}) dx_j = \pi(x_j^* | x_{-j})$$

For instance,  $\kappa_j$  can be a Metropolis–Hastings kernel on the lower dimensional space  $\mathcal{X}_j \ni x_j$ .

(Short hand notation  $x_{-j} = (x_1, \dots, x_{j-1}, x_{j+1}, \dots, x_d)$ .)

9/19

## ex) Metropolis-within-Gibbs

**Target:**

$$\pi(x_1, x_2) \propto \underbrace{\tilde{\pi}(x_1, x_2) = \exp\left(-\frac{1}{2}(2x_1 + \sin(6.28x_1))^2\right)}_{\tilde{\pi}(x_1)} \underbrace{\mathcal{N}(x_2 | x_1^3, 0.1)}_{\pi(x_2 | x_1)}$$

**Gibbs sampler:**

**Set**  $x_1[1] = 0, x_2[1] = 0$

**for**  $m = 2, \dots, M$

    Draw  $x_1[m] \sim \kappa_1(x[m-1], x_1^*)$ ;

    Draw  $x_2[m] \sim \pi(x_2 | x_1[m])$ .

where  $\kappa_1$  is a Metropolis–Hastings kernel for  $\pi(x_1 | x_2)$ .

Note that  $\pi(x_1 | x_2) = \frac{\pi(x_1, x_2)}{\pi(x_2)}$ . Hence, **conditionally on**  $x_2$ ,

$$\pi(x_1 | x_2) \propto \pi(x_1, x_2) \propto \tilde{\pi}(x_1, x_2).$$

10/19

## ex) Metropolis-within-Gibbs

**Algorithm 1** Metropolis-within-Gibbs sampler for toy problem

1. **Initialize:** Set  $x_1[1] = 0$ ,  $x_2[1] = 0$ .

2. **For**  $m = 2$  **to**  $M$ , **iterate:**

a. Sample  $x'_1 \sim \mathcal{N}(x_1 | x_1[m-1], 0.5^2)$ .

b. Sample  $u \sim \mathcal{U}[0, 1]$ .

c. Compute the acceptance probability

$$\alpha = \min \left( 1, \frac{\tilde{\pi}(x'_1, x_2[m-1])}{\tilde{\pi}(x_1[m-1], x_2[m-1])} \right)$$

d. Set

$$x_1[m] = \begin{cases} x'_1 & \text{if } u \leq \alpha \\ x_1[m-1] & \text{otherwise} \end{cases}$$

e. Draw  $x_2[m] \sim \pi(x_2 | x_1[m])$ .

11/19

## Gibbs sampling for dynamical systems

## ex) Gibbs sampling for linear Gaussian system

Simple LG-SSM,

$$\begin{aligned} X_t &= 0.9X_{t-1} + V_t, & V_t &\sim \mathcal{N}(0, \Theta_1), \\ Y_t &= X_t + E_t, & E_t &\sim \mathcal{N}(0, \Theta_2), \end{aligned}$$

With inverse-Gamma priors:  $\Theta_1 \sim \mathcal{IG}(0.1, 0.1)$ ,  $\Theta_2 \sim \mathcal{IG}(0.1, 0.1)$ .

**Task:** Compute  $p(\theta | y_{1:T})$  for a batch of  $T = 100$  observations.

The inverse-Gamma distribution is **conjugate prior** for an unknown variance of a Gaussian likelihood  $\Rightarrow$

$$\begin{aligned} p(\theta_1 | x_{0:T}, y_{1:T}) &= \mathcal{IG} \left( \theta_1 | 0.1 + \frac{T}{2}, 0.1 + \frac{1}{2} \sum_{t=1}^T (x_t - 0.9x_{t-1})^2 \right), \\ p(\theta_2 | x_{0:T}, y_{1:T}) &= \mathcal{IG} \left( \theta_2 | 0.1 + \frac{T}{2}, 0.1 + \frac{1}{2} \sum_{t=1}^T (y_t - x_t)^2 \right). \end{aligned}$$

12/19

## ex) Gibbs sampling for linear Gaussian system

**Gibbs sampler:**

**Initialize**  $\theta_1[1] = \theta_2[1] = 5$  (arbitrary!)


**for**  $m = 2, \dots, M$

- Draw  $x_{0:T}[m] \sim p(x_{0:T} | \theta[m-1], y_{1:T})$ ,  
by using Kalman smoothing techniques.
- Draw  $\theta[m] \sim p(\theta | x_{0:T}[m], y_{1:T})$ ,  
i.e., simulate  $\theta_1[m]$  and  $\theta_2[m]$  from their inverse-Gamma posteriors.

13/19

## ex) Gibbs sampling for linear Gaussian system


First 500 iterations of the Gibbs sampler for  $\theta_1$  and  $\theta_2$ .



14/19

## ex) Gibbs sampling for linear Gaussian system

Marginal posterior distributions,  $p(\theta_1 | y_{1:T})$  and  $p(\theta_2 | y_{1:T})$ , based on 50 000 iterations of the Gibbs sampler.



15/19

## Gibbs sampling for nonlinear dynamical systems

What about a general nonlinear/non-Gaussian dynamical system?

$$X_t | (X_{t-1} = x_{t-1}, \Theta = \theta) \sim p(x_t | x_{t-1}, \theta),$$

$$Y_t | (X_t = x_t, \Theta = \theta) \sim p(y_t | x_t, \theta),$$

$$X_0 \sim p(x_0), \quad \Theta \sim p(\theta).$$

**Gibbs sampler:**


- Draw  $\theta^* \sim p(\theta | x_{0:T}, y_{1:T})$ , **OK!**
- Draw  $x_{0:T}^* \sim p(x_{0:T} | \theta^*, y_{1:T})$ . **Hard!**

**Problem:**  $p(x_{0:T} | \theta, y_{1:T})$  not available!

**Idea:** Approximate  $p(x_{0:T} | \theta, y_{1:T})$  using a particle filter?

16/19

## Sampling based on the PF



With  $\mathbb{P}(X_{0:T}^* = x_{0:T}^i) = w_T^i$  we get  $X_{0:T}^* \stackrel{\text{approx.}}{\sim} p(x_{0:T} | \theta, y_{1:T})$ .

17/19

Problems with this approach:

- Based on a PF  $\Rightarrow$  approximate sample.
- $p(\theta, x_{1:T} | y_{1:T})$  is not a stationary distribution.
- Relies on large  $N$  to be successful.
- A lot of wasted computations.

The PMCMC framework allows us to address these issues!

**Gibbs sampler:** an MCMC sampler that iteratively simulates the unknown variables of the model from their conditional distributions.

**MCMC within Gibbs:** If exact sampling from some conditional is not possible, we may use any valid MCMC kernel within a Gibbs sampler to simulate from this conditional.

**Gibbs sampling for dynamical systems:** boils down to sampling the model parameters **with fixed states** + sampling the states with **fixed parameters** (state inference).