



UPPSALA
UNIVERSITET

Sequential Monte Carlo Methods

Lecture 17 – SMC for Probabilistic Programs

Lawrence Murray, Uppsala University

2017-08-29

Aim:

- Introduce probabilistic programming as a modeling paradigm.
- Demonstrate SMC as an appropriate inference method.

Outline:

1. Probabilistic programs.
2. Some examples in Birch.
3. SMC for probabilistic programs.

Probabilistic programs

Programs as stochastic processes

- Consider a program that depends on random numbers.
- Execute that program on a processor.
- As it runs, its memory state evolves **dynamically** and **stochastically** in time.

Programs as stochastic processes

- Consider a program that depends on random numbers.
- Execute that program on a processor.
- As it runs, its memory state evolves **dynamically** and **stochastically** in time.

We can think of the running program as a **stochastic process**.

Programs as stochastic processes

- Let $k = 1, 2, \dots$ denote a sequence of **checkpoints**.
- Let $(x_{1:k})_{k \geq 1}$ denote the (memory) state of the running program at checkpoint k , where $x_{1:k} \in \mathcal{X}_{1:k}$ and $\mathcal{X}_{1:k} = \mathcal{X}_k \times \mathcal{X}_{1:k-1}$.
- The state transitions according to $p_k(x_k | x_{1:k-1})$.

Programs as stochastic processes

- Let $k = 1, 2, \dots$ denote a sequence of **checkpoints**.
- Let $(x_{1:k})_{k \geq 1}$ denote the (memory) state of the running program at checkpoint k , where $x_{1:k} \in \mathcal{X}_{1:k}$ and $\mathcal{X}_{1:k} = \mathcal{X}_k \times \mathcal{X}_{1:k-1}$.
- The state transitions according to $p_k(x_k | x_{1:k-1})$.

At each checkpoint we can manipulate the running program: pause execution, inspect memory state, consider distributions over that state, modify that state. This is what facilitates inference.

What is probabilistic programming?

Probabilistic programming is a programming paradigm that emphasises this perspective on programs.

Consider other programming paradigms that emphasise other perspectives: functional, imperative, object-oriented, aspect-oriented.

What is a probabilistic programming language?

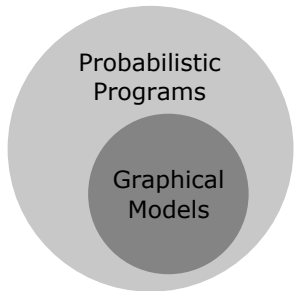
A **probabilistic programming language** (PPL) is a programming language that provides ergonomic support for the probabilistic programming paradigm.

A PPL may provide, for example:

- A library of probability distributions with the ability to evaluate and simulate them.
- Specialised language features for specifying probabilistic models.
- Specialised language features for writing probabilistic inference methods.

What is a probabilistic program?

A **probabilistic program** encodes a **probabilistic model** according to the semantics of a particular **probabilistic programming language**.



Probabilistic programs extend graphical models with support for **stochastic branching**.

The particular PPL that we will use is **Birch**, which is currently being developed at Uppsala University.

- It is the successor of **LibBi** (www.libbi.org).
- It is a probabilistic and object-oriented language.
- It compiles down to C++.

Example #1

$x \sim \text{Gaussian}(0.0, 1.0);$

$y \sim \text{Gaussian}(x, 1.0);$

$z \sim \text{Gaussian}(y, 1.0);$

Example #1

```
x ~ Gaussian(0.0, 1.0);  
y ~ Gaussian(x, 1.0);  
z ~ Gaussian(y, 1.0);
```

Adopting **operational semantics**, the interpretation of a program is defined by its execution. Here, the program encodes a **joint distribution**.

Example #1

$x \sim \text{Gaussian}(0.0, 1.0);$

$y \sim \text{Gaussian}(x, 1.0);$

$z \sim \text{Gaussian}(y, 1.0);$

$p(x)$



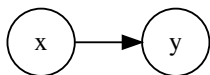
Example #1

$x \sim \text{Gaussian}(0.0, 1.0);$

$y \sim \text{Gaussian}(x, 1.0);$

$z \sim \text{Gaussian}(y, 1.0);$

$p(y|x)$



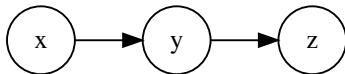
Example #1

$x \sim \text{Gaussian}(0.0, 1.0);$

$y \sim \text{Gaussian}(x, 1.0);$

$z \sim \text{Gaussian}(y, 1.0);$

$p(z|y)$



Example #1

$x \sim \text{Gaussian}(0.0, 1.0);$

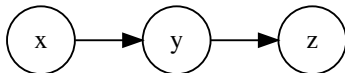
$y \sim \text{Gaussian}(x, 1.0);$

$z \sim \text{Gaussian}(y, 1.0);$

$p(x)$

$p(y|x)$

$p(z|y)$



Example #2

```
 $\beta \sim \text{Bernoulli}(0.5);$   
 $x \sim \text{Gaussian}(0.0, 1.0);$   
if ( $\beta$ ) {  
   $y \sim \text{Gaussian}(x, 1.0);$   
} else {  
   $y \sim \text{Gaussian}(0.0, 1.0);$   
}
```

Example #2

```
 $\beta \sim \text{Bernoulli}(0.5);$   
 $x \sim \text{Gaussian}(0.0, 1.0);$   
if ( $\beta$ ) {  
   $y \sim \text{Gaussian}(x, 1.0);$   
} else {  
   $y \sim \text{Gaussian}(0.0, 1.0);$   
}
```

Example #2

```
 $\beta \sim \text{Bernoulli}(0.5);$   
 $x \sim \text{Gaussian}(0.0, 1.0);$   
if ( $\beta$ ) {  
   $y \sim \text{Gaussian}(x, 1.0);$   
} else {  
   $y \sim \text{Gaussian}(0.0, 1.0);$   
}
```

$p(\beta)$



Example #2

```
 $\beta \sim \text{Bernoulli}(0.5);$   
 $x \sim \text{Gaussian}(0.0, 1.0);$   
if ( $\beta$ ) {  
   $y \sim \text{Gaussian}(x, 1.0);$   
} else {  
   $y \sim \text{Gaussian}(0.0, 1.0);$   
}
```

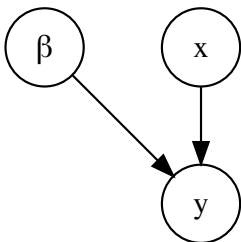
$p(x)$



Example #2

```
 $\beta \sim \text{Bernoulli}(0.5);$   
 $x \sim \text{Gaussian}(0.0, 1.0);$   
if ( $\beta$ ) {  
   $y \sim \text{Gaussian}(x, 1.0);$   
} else {  
   $y \sim \text{Gaussian}(0.0, 1.0);$   
}
```

$$p(y|x, \beta)$$



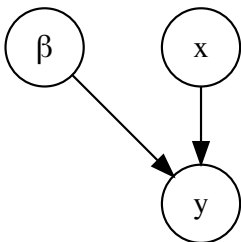
Example #2

```
 $\beta \sim \text{Bernoulli}(0.5);$   
 $x \sim \text{Gaussian}(0.0, 1.0);$   
if ( $\beta$ ) {  
   $y \sim \text{Gaussian}(x, 1.0);$   
} else {  
   $y \sim \text{Gaussian}(0.0, 1.0);$   
}
```

$p(\beta)$

$p(x)$

$p(y|x, \beta)$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

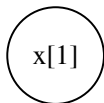

Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

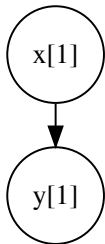
$p(x_1)$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

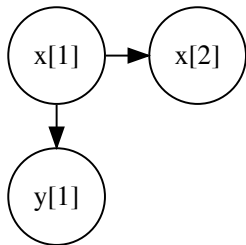
$$p(y_1 | x_1)$$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

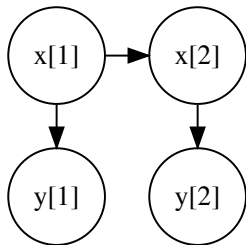
$$p(x_t | x_{t-1})$$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

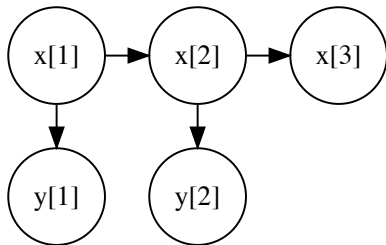
$$p(y_t | x_t)$$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

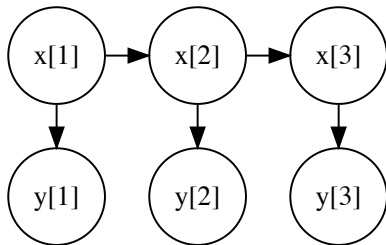
$$p(x_t | x_{t-1})$$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

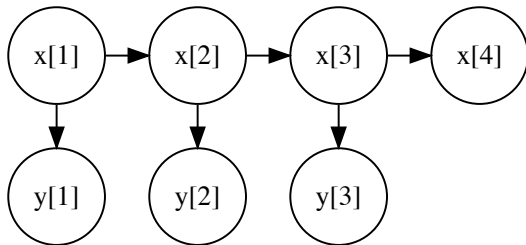
$$p(y_t | x_t)$$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

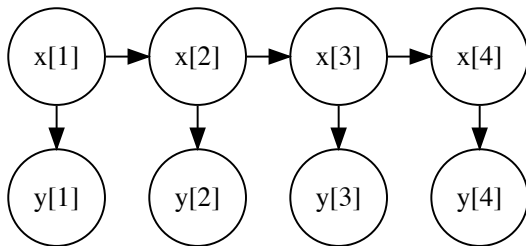
$$p(x_t | x_{t-1})$$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

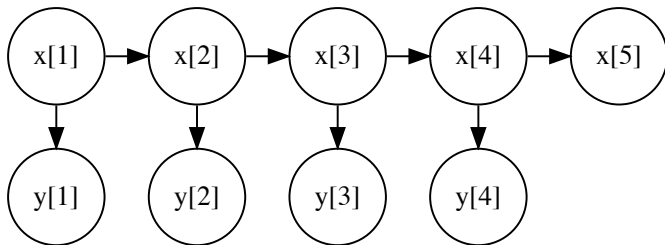
$$p(y_t | x_t)$$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

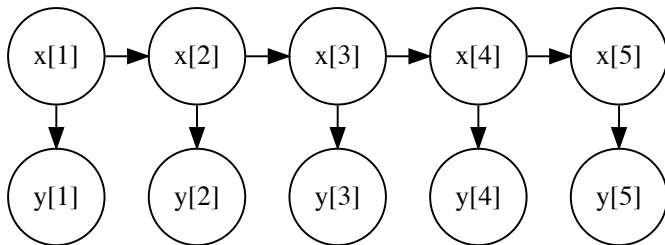
$p(x_t | x_{t-1})$



Example #3

```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

$$p(y_t | x_t)$$



Example #3

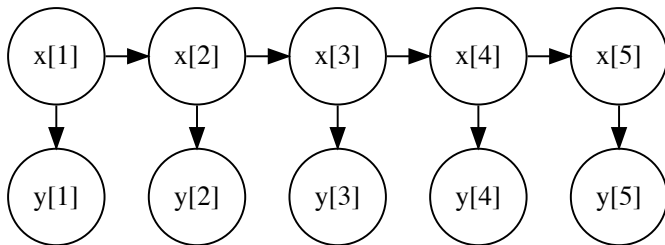
```
x[1] ~ Gaussian(0.0, 1.0);  
y[1] ~ Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] ~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~ Gaussian(x[t], 1.0);  
}
```

$$p(x_1)$$

$$p(y_1 | x_1)$$

$$p(x_t | x_{t-1})$$

$$p(y_t | x_t)$$



Checkpoints

The PPL will define checkpoints when interesting events happen in the running of the program. A typical setup uses two categories of checkpoint:

1. **Sample** when x is distributed according to some distribution p and should be sampled.
2. **Observe** when x is distributed according to some distribution p and should be observed to have some given value.

In Birch, these are triggered by special operators:

1. $x \leftarrow p$
2. $x \rightsquigarrow p$

Example (Checkpoints)

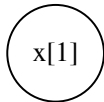
```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

Now, the program explicitly states which variables must be sampled, and which have given values and should be observed. The program encodes a **posterior distribution**.

Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

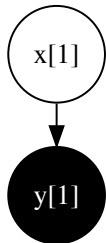
sample(x[1])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

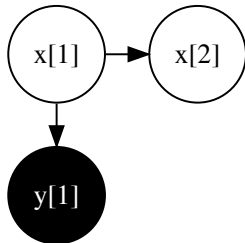
observe(x[1])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

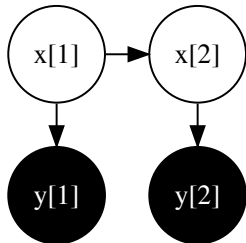
sample(x[t])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

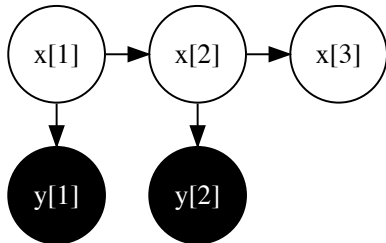
observe(x[t])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

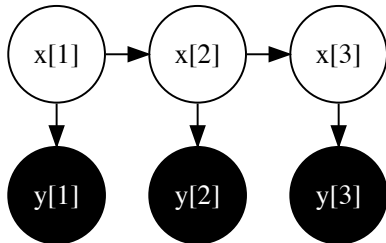
sample(x[t])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

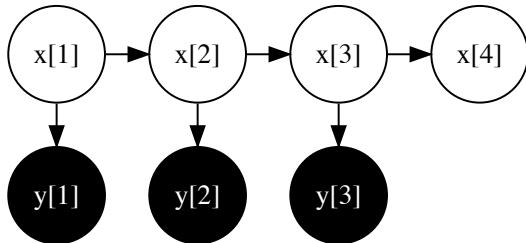
observe(x[t])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

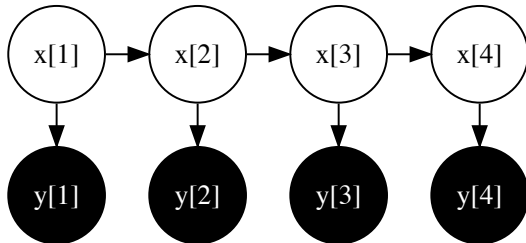
sample(x[t])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

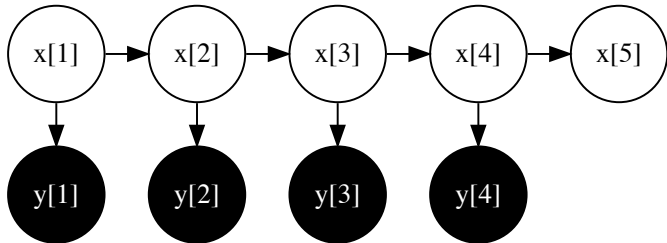
observe(x[t])



Example (Checkpoints)

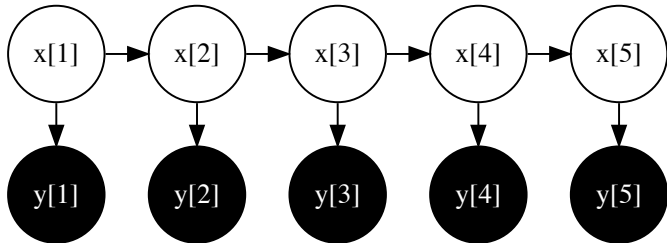
```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

sample(x[t])



Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
  observe(x[t])  
}
```

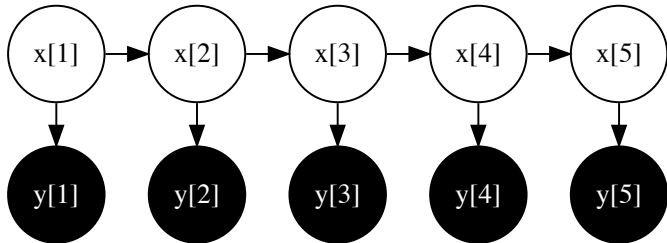


Example (Checkpoints)

```
x[1] <~ Gaussian(0.0, 1.0);  
y[1] ~> Gaussian(x[1], 1.0);  
for (t in 2..T) {  
  x[t] <~ Gaussian(a*x[t - 1], 1.0);  
  y[t] ~> Gaussian(x[t], 1.0);  
}
```

sample(x[1])
observe(x[1])

sample(x[t])
observe(x[t])



A probabilistic program encodes a **probabilistic model**.

A running probabilistic program is a **stochastic process**.

SMC for probabilistic programs

SMC for probabilistic programs

Recall that SMC can be used to approximate a sequence of probability distributions on a sequence of probability spaces of increasing dimension.

Let $\{\pi_k(\mathbf{x}_{1:k})\}_{k \geq 1}$ be the sequence of target distributions

$$\pi_k(\mathbf{x}_{1:k}) = \frac{\tilde{\pi}_k(\mathbf{x}_{1:k})}{Z_k}$$

Where

$$\pi_k(\mathbf{x}_{1:k}) \approx \sum_{i=1}^N w_k^i \delta_{\mathbf{x}_{1:k}^i}(\mathbf{x}_{1:k})$$

and the weighted particle populations $\{\mathbf{x}_{1:k}^i, w_k^i\}_{i=1}^N$ are generated sequentially for $k = 1, 2, \dots$

We can use SMC for inference on running probabilistic programs.

- Each of the N particles is a running probabilistic program.
- We have:

$$\begin{aligned}\tilde{\pi}_k(x_{1:k}) &= p_k(x_k \mid x_{1:k-1})\tilde{\pi}_{k-1}(x_{1:k-1}) \\ q_k(x_k \mid x_{1:k-1}) &= p_k(x_k \mid x_{1:k-1}).\end{aligned}$$

- That is, the probabilistic program defines the target and the proposal, much like the bootstrap particle filter.

Assume that we have obtained $\{x_{1:k-1}^i, w_{k-1}^i\}_{i=1}^N$.

1. **Resample**: Sample a_k^i with $\mathbb{P}(a_k^i = j) = \nu_{k-1}^j, j = 1, \dots, N$.
2. If this is a **sample** checkpoint, then **propagate**:

$$x_k^i \sim p_k(x_k | x_{1:k-1}^{a_k^i}) \text{ and } x_{1:k}^i = (x_{1:k-1}^{a_k^i}, x_k^i)$$

3. If this is an **observe** checkpoint, then **weight**:

$$w_k^i \propto \frac{w_{k-1}^{a_k^i} p_k(x_k^i | x_{1:k-1}^{a_k^i})}{\nu_{k-1}^{a_k^i}}.$$

The result is a new weighted set of particles $\{x_{1:k}^i, w_k^i\}_{i=1}^N$.

A few concepts to summarize lecture 17

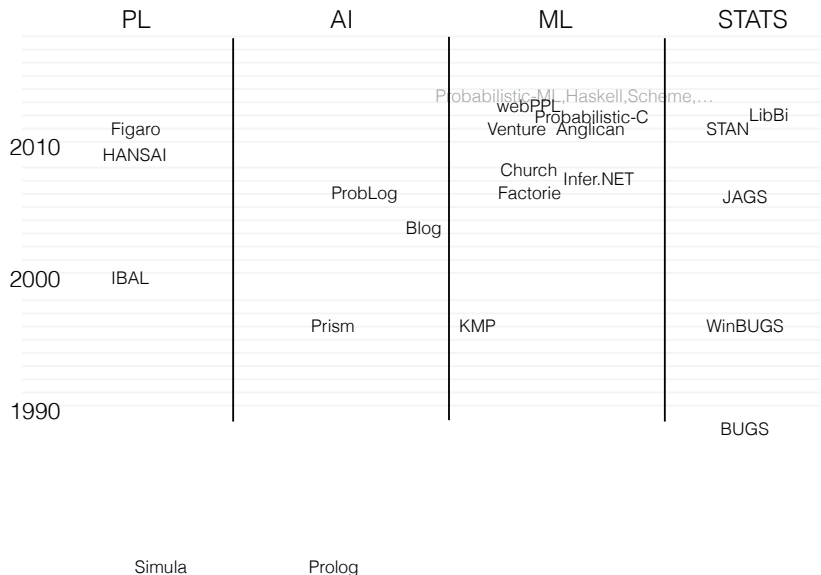
A **probabilistic program** encodes a **probabilistic model** according to the semantics of a particular **probabilistic programming language**.

The memory state of a running probabilistic program evolves dynamically and stochastically in time and so is a **stochastic process**.

General Sequential Monte Carlo can be applied to perform inference across a sequence of target distributions defined by the **checkpoints** of the running program.

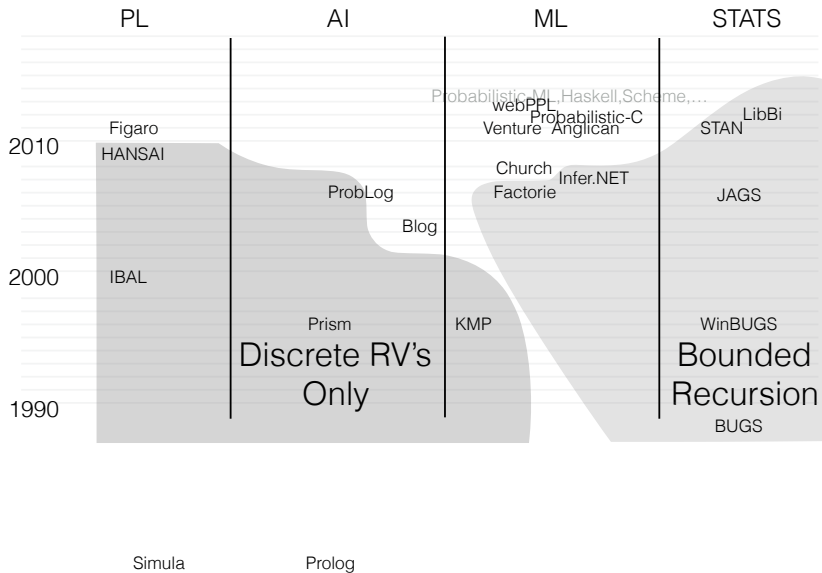
Further study

Probabilistic Programming Languages



(Figure courtesy of Frank Wood, Oxford.)

Probabilistic Programming Languages



(Figure courtesy of Frank Wood, Oxford.)